

2.5 Object detection

A pretrained maskrcnn_resnet50 model was used for this task. Since the images were different sizes, the training was not done in batches. Since the task used models that are a part of the PyTorch library the hardest part was figuring out how to convert the region of balloons to masks. I implemented it using the library skimage.draw. Here are some images that show which pixels got classified as balloons with more than 50% certainty but turning the rest of image to grayscale:



